

# Class 7 BGGN213

AUTHOR

Lisanne Stouthart (PID A69036187)

## Quarto

---

Quarto enables you to weave together content and executable code into a finished document. To learn more about Quarto see <https://quarto.org>.

## Class - Machine Learning 1

---

Before we get into clustering methods let's make some sample data to cluster where we know what the answer should be.

To help with this I will use the 'rnorm()' function.

`dnorm(x, mean = 0, sd = 1, log = FALSE)` `pnorm(q, mean = 0, sd = 1, lower.tail = TRUE, log.p = FALSE)`  
`qnorm(p, mean = 0, sd = 1, lower.tail = TRUE, log.p = FALSE)` `rnorm(n, mean = 0, sd = 1)`

```
# rnorm()  
# it has 3 arguments, where 2 out of the 3 arguments are fixed numbers.  
  
rnorm(5)
```


```
[1] -0.04001234  0.68597051  1.07432277 -0.37158666  0.37560176
```

```
rnorm(15)
```

```
[1] -0.83081203  0.36672117  1.52431777 -0.38404967 -1.27266604 -0.99357998  
[7] -0.40557806  0.01959658  1.48226903  0.80082539 -0.33246431  0.76634480  
[13]  2.01612289  0.21182065 -0.99710142
```


```
hist(rnorm(15000))
```

## Histogram of rnorm(150000)




```
hist(rnorm(150000, mean=3))
```

## Histogram of `rnorm(150000, mean = 3)`




```
hist(rnorm(150000, mean=-3))
```

## Histogram of rnorm(150000, mean = -3)



```
n = 10000  
hist(c(rnorm(n, mean=3), rnorm(n, mean=-3)))
```

## Histogram of `c(rnorm(n, mean = 3), rnorm(n, mean = -3))`



```
n = 30
c(rnorm(n, mean=3), rnorm(n, mean=-3))
```

```
[1] 2.3875991 2.1519626 5.1819411 2.8475886 2.9156329 2.5451940
[7] 1.3406880 2.7783780 4.6916243 3.6794071 2.2605492 3.8451309
[13] 1.2771424 2.8129824 2.4945789 1.5244540 4.2013811 0.9510227
[19] 4.7722170 3.6147070 3.1691506 2.2369681 3.5414158 2.2059774
[25] 2.6149123 3.5499059 2.1242428 3.1467887 2.5193566 4.2848344
[31] -4.4012033 -4.2292934 -3.2462846 -1.8146888 -4.3372222 -3.1793014
[37] -1.8153119 -3.6881381 -3.4270941 -3.6013313 -3.8240825 -4.6392731
[43] -2.2353999 -2.9303139 -3.2010897 -1.9443771 -5.0248378 -4.4468318
[49] -2.3281949 -2.6903217 -3.1669577 -2.6235725 -3.3530595 -2.2105007
[55] -2.8355858 -4.2625628 -4.0483448 -3.2362184 -3.5232899 -2.2982812
```

```
x <- c(rnorm(n, mean=3), rnorm(n, mean=-3))
x
```

```
[1] 2.4485957 2.8970112 3.6338778 2.9001816 2.9080854 3.6498481
[7] 2.9218679 3.8375588 3.8146197 2.1855044 4.2121325 3.9143063
[13] 2.0037811 3.2844479 2.3215824 4.0638247 4.3940141 1.5648858
[19] 2.3756995 4.6339226 3.1604437 2.3822958 0.2377887 3.0933090
[25] 4.1708672 2.0677726 2.1395113 1.2399405 3.5271081 2.0080333
```

```
[31] -3.3041363 -2.8431167 -2.5284681 -4.1403048 -3.8434953 -2.7201755
[37] -1.7230558 -4.1083153 -3.3165389 -2.0127575 -4.3233424 -3.6131621
[43] -3.8747542 -2.3163600 -1.3576175 -2.1894470 -3.4805658 -1.0772638
[49] -1.2172450 -2.1500511 -2.5551595 -5.4395708 -2.1455182 -2.3519747
[55] -2.7612567 -1.8270459 -2.5132285 -5.6599396 -2.7580815 -2.5259105
```

```
y <- rev(x)
y
```

```
[1] -2.5259105 -2.7580815 -5.6599396 -2.5132285 -1.8270459 -2.7612567
[7] -2.3519747 -2.1455182 -5.4395708 -2.5551595 -2.1500511 -1.2172450
[13] -1.0772638 -3.4805658 -2.1894470 -1.3576175 -2.3163600 -3.8747542
[19] -3.6131621 -4.3233424 -2.0127575 -3.3165389 -4.1083153 -1.7230558
[25] -2.7201755 -3.8434953 -4.1403048 -2.5284681 -2.8431167 -3.3041363
[31] 2.0080333 3.5271081 1.2399405 2.1395113 2.0677726 4.1708672
[37] 3.0933090 0.2377887 2.3822958 3.1604437 4.6339226 2.3756995
[43] 1.5648858 4.3940141 4.0638247 2.3215824 3.2844479 2.0037811
[49] 3.9143063 4.2121325 2.1855044 3.8146197 3.8375588 2.9218679
[55] 3.6498481 2.9080854 2.9001816 3.6338778 2.8970112 2.4485957
```


```
z <- cbind(x,y)
z
```

```
      x      y
[1,] 2.4485957 -2.5259105
[2,] 2.8970112 -2.7580815
[3,] 3.6338778 -5.6599396
[4,] 2.9001816 -2.5132285
[5,] 2.9080854 -1.8270459
[6,] 3.6498481 -2.7612567
[7,] 2.9218679 -2.3519747
[8,] 3.8375588 -2.1455182
[9,] 3.8146197 -5.4395708
[10,] 2.1855044 -2.5551595
[11,] 4.2121325 -2.1500511
[12,] 3.9143063 -1.2172450
[13,] 2.0037811 -1.0772638
[14,] 3.2844479 -3.4805658
[15,] 2.3215824 -2.1894470
[16,] 4.0638247 -1.3576175
[17,] 4.3940141 -2.3163600
[18,] 1.5648858 -3.8747542
[19,] 2.3756995 -3.6131621
[20,] 4.6339226 -4.3233424
[21,] 3.1604437 -2.0127575
[22,] 2.3822958 -3.3165389
[23,] 0.2377887 -4.1083153
[24,] 3.0933090 -1.7230558
[25,] 4.1708672 -2.7201755
```

```
[26,] 2.0677726 -3.8434953
[27,] 2.1395113 -4.1403048
[28,] 1.2399405 -2.5284681
[29,] 3.5271081 -2.8431167
[30,] 2.0080333 -3.3041363
[31,] -3.3041363 2.0080333
[32,] -2.8431167 3.5271081
[33,] -2.5284681 1.2399405
[34,] -4.1403048 2.1395113
[35,] -3.8434953 2.0677726
[36,] -2.7201755 4.1708672
[37,] -1.7230558 3.0933090
[38,] -4.1083153 0.2377887
[39,] -3.3165389 2.3822958
[40,] -2.0127575 3.1604437
[41,] -4.3233424 4.6339226
[42,] -3.6131621 2.3756995
[43,] -3.8747542 1.5648858
[44,] -2.3163600 4.3940141
[45,] -1.3576175 4.0638247
[46,] -2.1894470 2.3215824
[47,] -3.4805658 3.2844479
[48,] -1.0772638 2.0037811
[49,] -1.2172450 3.9143063
[50,] -2.1500511 4.2121325
[51,] -2.5551595 2.1855044
[52,] -5.4395708 3.8146197
[53,] -2.1455182 3.8375588
[54,] -2.3519747 2.9218679
[55,] -2.7612567 3.6498481
[56,] -1.8270459 2.9080854
[57,] -2.5132285 2.9001816
[58,] -5.6599396 3.6338778
[59,] -2.7580815 2.8970112
[60,] -2.5259105 2.4485957
```

```
# z <- rbind(x,y)
# z
```

```
plot(z)
```



The function in base R (meaning that you dont have to install this) for k-means clustering is called 'kmeans()'.

```
km <- kmeans(z, center = 2)
# is 30,30 because we set 30 points each dataset
# he has slightly different answers but that is because it is random
km
```


Cluster means:

Clustering vector:


localhost:7396








```
plot(z, col=c("red", "blue"))
```




```
#this is random assigning of the blue and red, alternating
```

```
plot(z, col = 3) #green, color by number
```




```
plot(z, col=c(1,2)) #green, color by number, random again
```




```
plot(z, col=km$cluster) #plot with clustering result
```

```
plot(z, col=km$cluster) +  
points(km$centers, col="blue") #plot with clustering result and center result
```



`integer(0)`

```
plot(z, col=km$cluster) +  
points(km$centers, col="blue", pch=16, cex=3) #to make center bit bigger
```




```
integer(0)
```

```
# if you play with number of pch, you can change the shape  
# if you play with number of cex, you can change the size
```

## Can you cluster our data in 'z' into four clusters please?


```
km4 <- kmeans(z, center = 4)  
plot(z, col=km4$cluster) +  
  points(km4$centers, col="blue", pch=16, cex=3)
```



integer(0)

```
km5 <- kmeans(z, center = 5)
plot(z, col=km5$cluster) +
  points(km5$centers, col="blue", pch=15, cex=2)
```





```
integer(0)
```

## Hierarchical Clustering

Hierarchical cluster analysis on a set of dissimilarities and methods for analyzing it. `hclust(d, method = "complete", members = NULL)`

The main function for hierarchical clustering in base R is called '`hclust()`'. Unlike '`kmeans()`' I can not just pass in my data as input I first need a distance matrix from my data.

```
d <- dist(z)
hc <- hclust(d)
hc # this is not really helpful
```


Call:


```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

There is a specific `hclust plot()` method ...

## Cluster Dendrogram





## Principal Component Analysis

Principal component analysis (PCA) is a well established “multivariate statistical technique” used to reduce the dimensionality of a complex data set to a more manageable number (typically 2D or 3D). This method is particularly useful for highlighting strong patterns and relationships in large datasets (i.e. revealing major similarities and differences) that are otherwise hard to visualize. As we will see again and again in this course PCA is often used to make all sorts of bioinformatics data easy to explore and visualize.

### PCA of UK food data

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93

Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Q1: How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
## Complete the following code to find out how many rows and columns are in x?
nrow(x) #answer = 17 rows
```

[1] 17

```
ncol(x) #answer = 5 columns
```

[1] 4

```
dim(x) #both combined
```

[1] 17 4

```
## Preview the first 6 rows
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
# Note how the minus indexing works
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	Wales	Scotland	N.Ireland
105	103	103	66
245	227	242	267

685	803	750	586
147	160	122	93
193	235	184	209
156	175	147	139

```
# Check dimensions again
dim(x) # answer is 17 4
```

```
[1] 17 3
```

```
x <- read.csv(url, row.names=1)
head(x)
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2: Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?


```
# x <- read.csv(url, row.names=1)
# This is safe
```

Q3: Changing what optional argument in the above barplot() function results in the following plot?

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x))) #answer is change to F
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



```
# It always compares two countries. Column 2, row 1, compares england with wales.
# If it is not on the line, it is different between the countries. If the point is
```

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

## PCA to the rescue

The main function to do PCA in base R is called 'prcomp()'.

```
# Use the prcomp() PCA function
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's see what is inside our result object 'pca' that we just calculated:



```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```

```
[1] "prcomp"
```

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

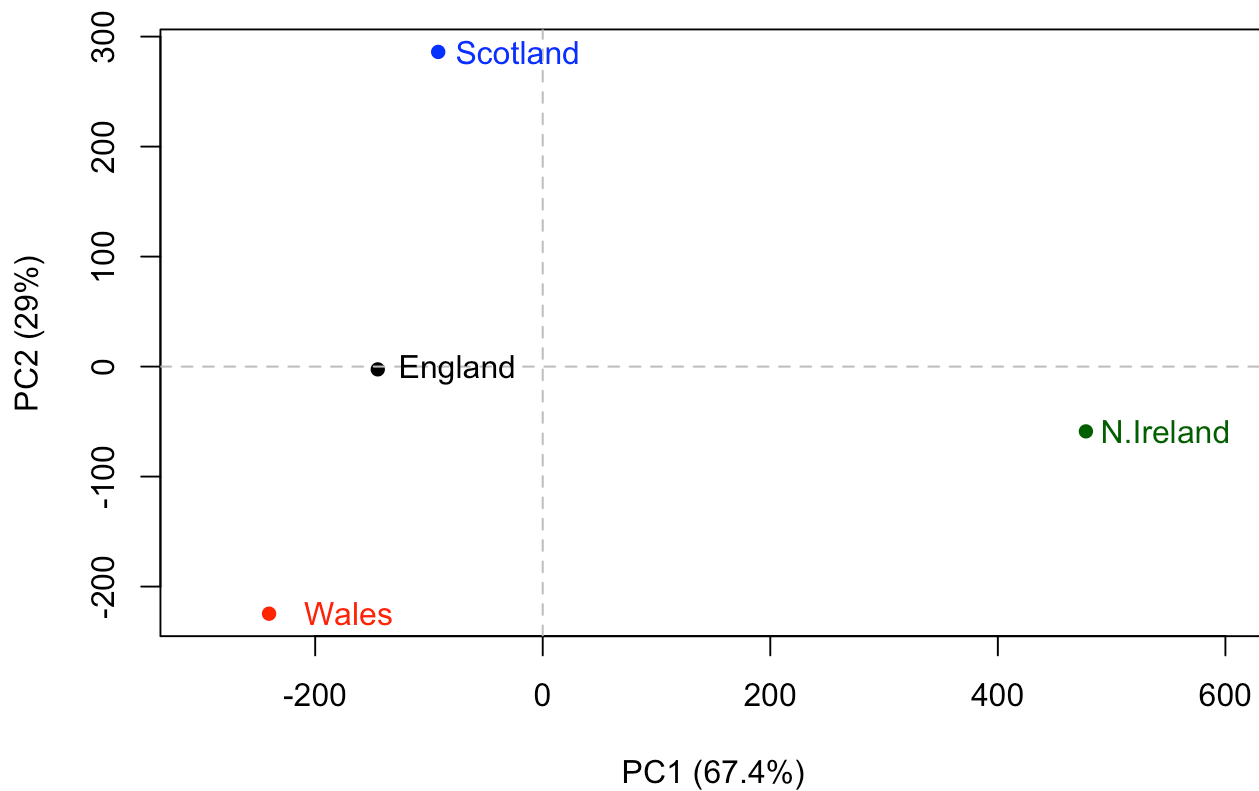
```
pca$x #if you want to data
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
# To make our main result figure, called a "PC plot" (or "score plot", or "ordina
```

```
# Plot PC1 vs PC2
```

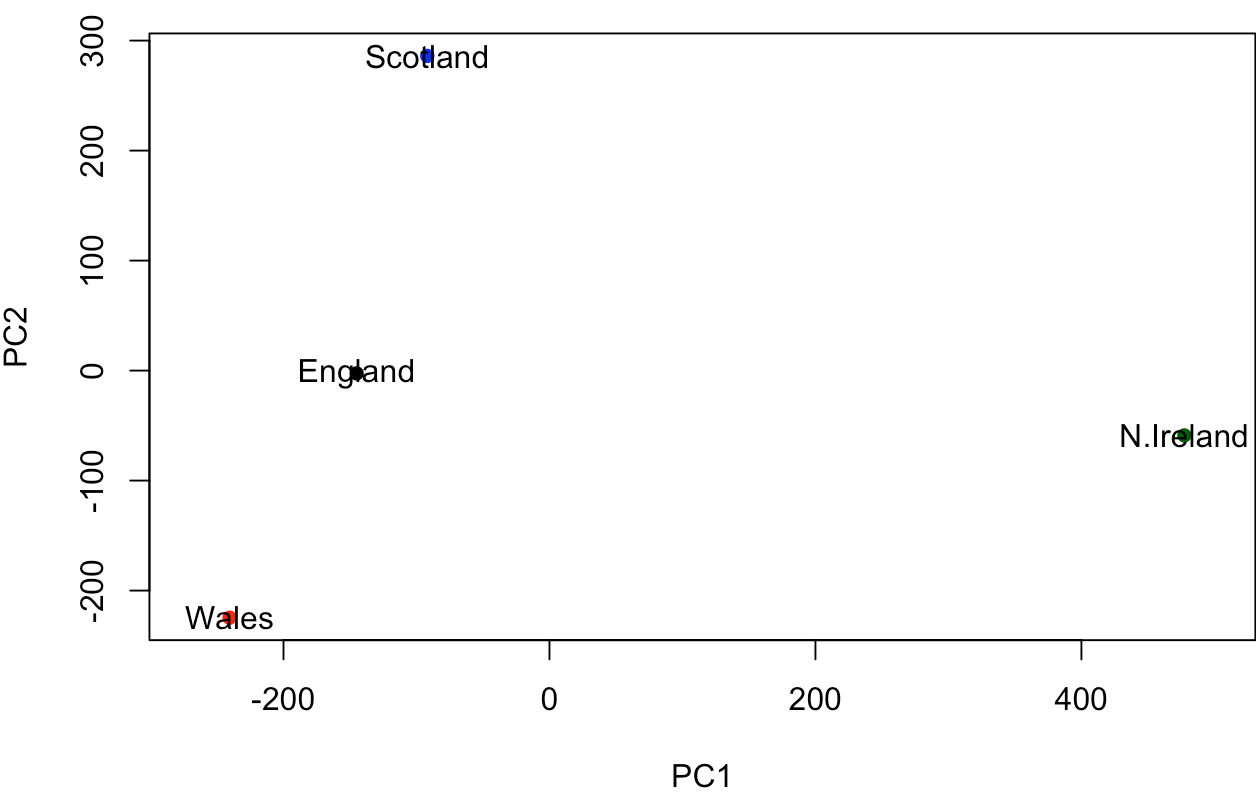
```
plot(pca$x[,1], pca$x[,2],
     col=c("black", "red", "blue", "darkgreen"),
     pch=16,
     xlab="PC1 (67.4%)", ylab="PC2 (29%)",
     xlim=c(-300,600)) +
abline(h=0, col="gray", lty=2) +
abline(v=0, col="gray", lty=2) +
text(pca$x[,1]+70, pca$x[,2], col=names(x),
     col=c("black", "red", "blue", "darkgreen"))
```



`integer(0)`

Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500), col=c("black", "blue", "red", "green"))
text(pca$x[,1], pca$x[,2], col=names(x))
```




Q9: Generate a similar ‘loadings plot’ for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.409382587
Carcass_meat	0.047927628	0.013915823	0.06367111	0.729481922
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.331001134
Fish	-0.084414983	-0.050754947	0.03906481	0.022375878
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.034512161
Sugars	-0.037620983	-0.043021699	-0.03605745	0.024943337
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	0.021396007
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	0.001606882
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.031153231
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	-0.017379680
Processed_Veg	-0.036488269	-0.045451802	0.05289191	0.021250980
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.227657348
Cereals	-0.047702858	-0.212599678	-0.35884921	0.100043319
Beverages	-0.026187756	-0.030560542	-0.04135860	-0.018382072
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.222319484
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.273126013
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001890737

```
## Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



## PCA of RNA-seq data

Q10: How many genes and samples are in this data set?